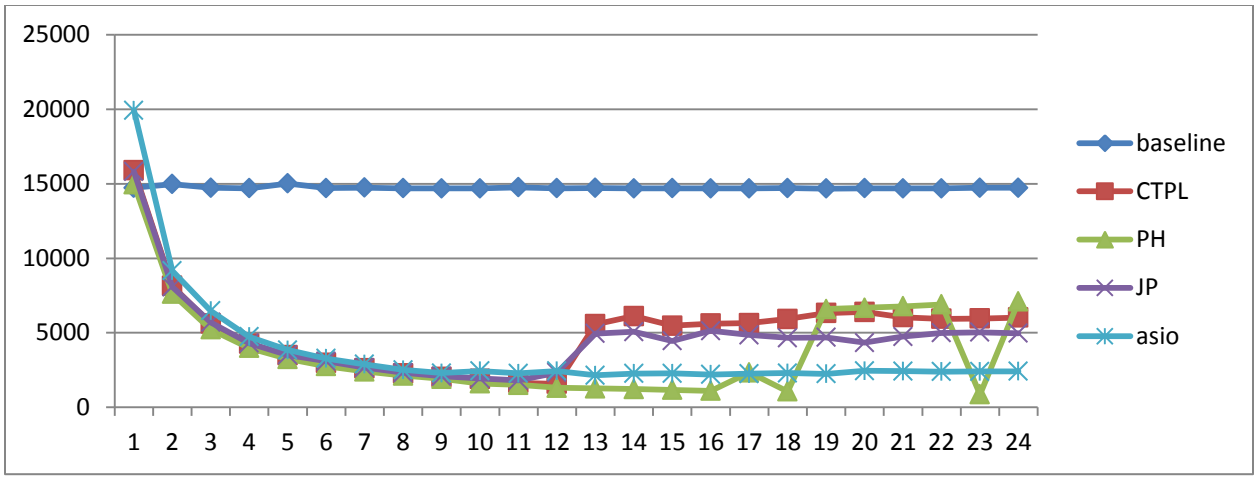

Execution time in nanoseconds, per iteration, as a function of the number of threads, of each of the tested thread pools:
Bassline – single thread
CTPL - https://github.com/vit-vit/CTPL
PH - https://github.com/henkel/threadpool
JP - https://github.com/progschj/ThreadPool
asio - http://stackoverflow.com/questions/12215395/thread-pool-using-boost-asio
No information presented for boost executor, since Centos7 (on which the test was done) does not have the correct boost version.

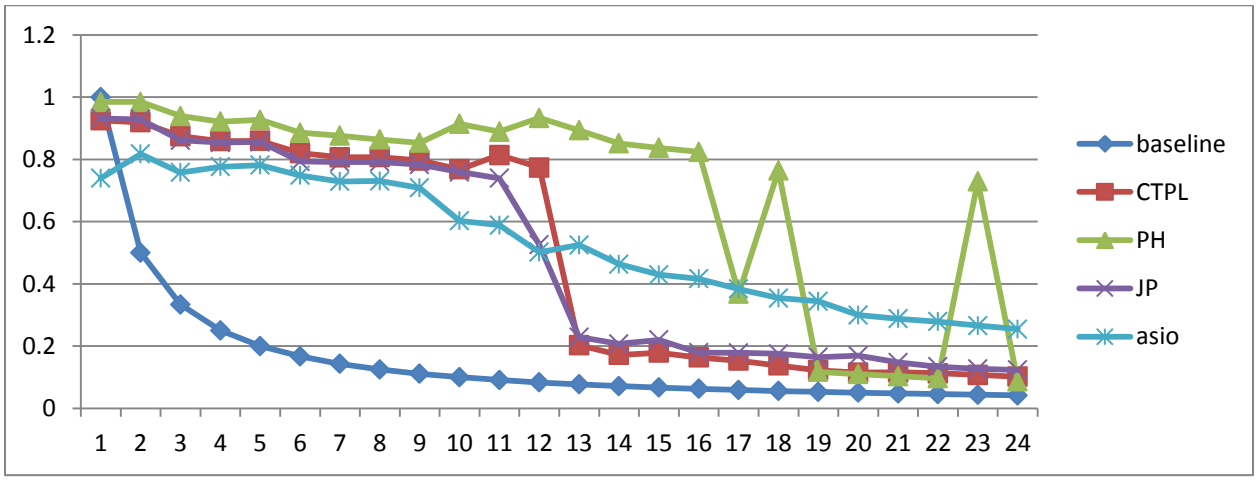

Efficiency ratio between baseline and the tested thread pool, divided by number of threads. Most implementations are “almost linear” (more than 80% sub-linear) up to ~12 threads, and then efficiency drops significantly.

Operating system: Centos7 kernel: 3.10.0-693.11.6.el7.x86_64
Hardware: 24 cores, Intel(R) Xeon(R) CPU X5680@3.33GHz